

ICS Homework 13

June 6, 2020

1 Organization

1.1

Let's consider about a page-removal algorithm: clock algorithm. Suppose we have a primary device which has 3 physical blocks, every time a reference string P come, it will follow the pseudo-code:

```
1  if hit(P)
2      res_block = block contains P
3      referenced_bit[res_block] <- True
4      clock_arm does not change
5  else
6      while referenced_bit[clock_arm]
7          referenced_bit[clock_arm] <- False
8          clock_arm <- next block
9      res_block = block[clock_arm]
10     referenced_bit[clock_arm] <- True
11     clock_arm <- next block
12 return res_block
```

Fill in the table (If you don't know what to fill just write down a '-'). Note: '*' means the position of the clock arm; you also need to tell what the referenced bit is for each page block at that time (1 for True and 0 for False).

Time	0	1	2	3	4	5	6	7	8
Reference string	-	3	4	2	6	4	3	7	4
Primary Device Contents	*	3	3	3 *	6	6	6 *	6	4
	-	*	4	4	4 *	4 *	4	7	7 *
	-	-	*	2	2	2	3	3 *	3
Referenced Bit	0	1	1	1	1	1	1	0	1
	0	0	1	1	0	1	0	1	1
	0	0	0	1	0	0	1	1	0
Page Absent	-	Y	Y	Y	Y	N	Y	Y	Y

1.2

Please consult the related information. Find and complement the definition of struct vm_area_struct in Linux v5.7.

(https://elixir.bootlin.com/linux/v5.7/source/include/linux/mm_types.h#L297)

```
1 struct vm_area_struct {
2     /* The first cache line has the info for VMA tree walking.
3      */
4     unsigned long vm_start; /* Our start address within
5                             vm_mm. */
6     unsigned long vm_end; /* The first byte after our end
7                             address
8                             within vm_mm. */
9     /* linked list of VM areas per task, sorted by address */
10    struct vm_area_struct *vm_next, *vm_prev;
11
12    struct rb_node vm_rb;
13
14    /*
15     * Largest free memory gap in bytes to the left of this VMA.
16     * Either between this VMA and vma->vm_prev, or
17     * between one of the
18     * VMAs below us in the VMA rbtree and its ->vm_prev.
19     * This helps
20     * get_unmapped_area find a free area of the right size.
21     */
22    unsigned long rb_subtree_gap;
23
24    /* Second cache line starts here. */
25
26    struct mm_struct *vm_mm; /* The address space we
27                             belong to. */
28
29    /*
30     * Access permissions of this VMA.
31     * See vmf_insert_mixed_prot() for discussion.
32     */
33    pgprot_t vm_page_prot;
34    unsigned long vm_flags; /* Flags, see mm.h. */
35
36    /*
37     * For areas with an address space and backing store,
38     * linkage into the address_space->i_mmap interval tree.
```

```

35     */
36     struct {
37         struct rb_node rb;
38         unsigned long rb_subtree_last;
39     } shared;
40
41     /*
42     * A file's MAP_PRIVATE vma can be in both i_mmap tree
43     * and anon_vma
44     * list, after a COW of one of the file pages. A
45     * MAP_SHARED vma
46     * can only be in the i_mmap tree. An anonymous
47     * MAP_PRIVATE, stack
48     * or brk vma (with NULL file) can only be in an anon_vma
49     * list.
50     */
51     struct list_head anon_vma_chain; /* Serialized by
52     mmap_sem &
53     * page_table_lock */
54     struct anon_vma *anon_vma; /* Serialized by
55     page_table_lock */
56
57     /* Function pointers to deal with this struct. */
58     const struct vm_operations_struct *vm_ops;
59
60     /* Information about our backing store: */
61     unsigned long vm_pgoff; /* Offset (within vm_file) in
62     PAGE_SIZE
63     units */
64     struct file * vm_file; /* File we map to (can be NULL
65     ). */
66     void * vm_private_data; /* was vm_pte (shared mem)
67     */
68
69 #ifdef CONFIG_SWAP
70     atomic_long_t swap_readahead_info;
71 #endif
72 #ifndef CONFIG_MMU
73     struct vm_region *vm_region; /* NOMMU mapping
74     region */
75 #endif
76 #ifdef CONFIG_NUMA
77     struct mempolicy *vm_policy; /* NUMA policy for
78     the VMA */
79 #endif
80     struct vm_userfaultfd_ctx vm_userfaultfd_ctx;

```

```
70 } __randomize_layout;
```

2 System Software

2.1

Let p denote the number of producers, c the number of consumers, and n the buffer size in units of items. Consider the following buffer implementation. For each of the following scenarios, indicate whether the **mutex** semaphore is necessary or not to implement function **sbuf_insert** and **sbuf_remove**.

```
1 typedef struct {
2     int *buf;      /* Buffer array */
3     int n;         /* Maximum number of slots */
4     int front;     /* buf[(front+1)%n] is first item */
5     int rear;      /* buf[rear%n] is last item */
6     sem_t mutex;   /* Protects accesses to buf */
7     sem_t slots;   /* Counts available slots */
8     sem_t items;   /* Counts available items */
9 } sbuf_t;
```

A. $p = 1, c = 1, n > 1$

B. $p = 1, c = 1, n = 1$

C. $p > 1, c > 1, n = 1$

ANS:

A. $p = 1, c = 1, n > 1$: No. Using semaphore **slots** and **items** is enough, and there is no concurrent access to **front** and **rear**.

B. $p = 1, c = 1, n = 1$: No. This case is the same as case A.

C. $p > 1, c > 1, n = 1$: No. In this case, **front** and **rear** are always 0. Consumers will be waiting for **items**, and producers will be waiting for **slots**.