

题号: V1

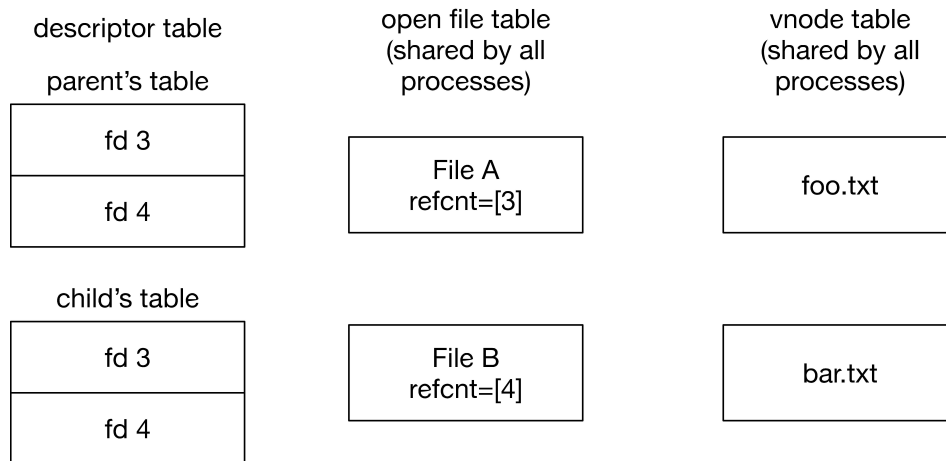
IO1 (16 points)

```
1. #include <csapp.h>
2.
3. int main(void) {
4.     int fd1, fd2;
5.     int pid;
6.     char c[4] = "ics";
7.     fd1 = Open("foo.txt", O_RDWR, 0);
8.     fd2 = Open("bar.txt", O_RDWR, 0);
9.     Write(fd2, c, 2);
10.    Dup2(1, fd2);
11.
12.    if (Fork() == 0) {
13.        Write(fd2, c, 1);
14.        Read(fd1, c, 1);
15.        Write(fd1, c, 2);
16.        Write(fd2, c, 2);
17.        Dup2(fd1, fd2);
18.    }
19.    else
20.        Wait(NULL);
21.
22.    Write(fd1, c, 2);
23.    Write(fd2, c, 2);
24.
25.    return 0;
26.}
```

NOTE: Initially, `foo.txt` contains "online-exam"; `bar.txt` contains "easy"; No error occurs in the execution. **NOTE:** suppose that read and write operations are atomic.

1. Please write down the output on **screen**. (2')
2. Right before the child process exit, please fill in the blanks in the **figure and table** below, like **Figure 10.12** in your text book. (NOTE: fd 0, 1, 2 are ignored, you need to complete **4 blanks**) (4')

descriptor table	open file table	vnode table
fd=3	File A	foo.txt
fd=4	[1]	[2]



- Please write down the **contents** of `foo.txt` and `bar.txt`. (6')
- If we change **line 13** from `write(fd2, c, 1)` to `printf("%c", c[0])`, write down the output on screen. (2') (**Hint:** `printf` has a buffer) If we want the same output as before, how to modify the code? (NOTE: you can't modify `printf`) (2')

题号: V2

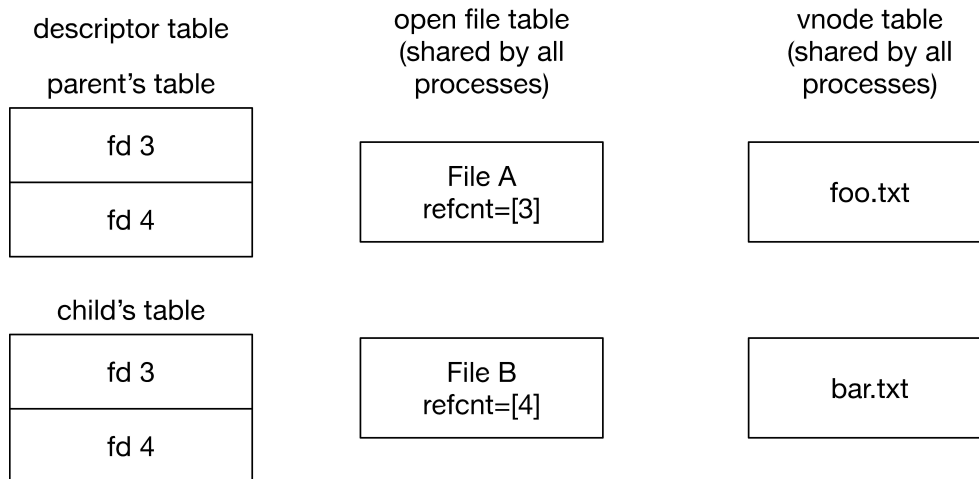
IO2 (16 points)

```
1. #include <csapp.h>
2.
3. int main(void) {
4.     int fd1, fd2;
5.     int pid;
6.     char c[4] = "NOP";
7.     fd1 = Open("foo.txt", O_RDWR, 0);
8.     fd2 = Open("bar.txt", O_RDWR, 0);
9.     Write(fd2, c, 2);
10.    Dup2(1, fd2);
11.
12.    if (Fork() == 0) {
13.        Write(fd2, c, 1);
14.        Read(fd1, c, 1);
15.        Write(fd1, c, 2);
16.        Write(fd2, c, 2);
17.        Dup2(fd1, fd2);
18.    }
19.    else
20.        Wait(NULL);
21.
22.    Write(fd1, c, 2);
23.    Write(fd2, c, 2);
24.
25.    return 0;
26.}
```

NOTE: Initially, `foo.txt` contains "ICS2020"; `bar.txt` contains "SJTU"; No error occurs in the execution. **NOTE:** suppose that read and write operations are atomic.

1. Please write down the output on **screen**. (2')
2. Right before the child process exit, please fill in the blanks in the **figure and table** below, like **Figure 10.12** in your text book. (NOTE: fd 0, 1, 2 are ignored, you need to complete **4 blanks**) (4')

descriptor table	open file table	vnode table
fd=3	File A	foo.txt
fd=4	[1]	[2]



3. Please write down the **contents** of **foo.txt** and **bar.txt**. (6')
4. If we change **line 13** from **write(fd2, c, 1)** to **printf("%c", c[0])**, write down the output on screen. (2') (**Hint**: **printf** has a buffer) If we want the same output as before, how to modify the code? (NOTE: you can't modify **printf**) (2')

题号: **W1**

Schedule (16 points)

We have following jobs in the workload. No I/O issues are involved.

Job	Arrival Time	Length of Run-time
A	0ms	8ms
B	4ms	4ms
C	6ms	6ms
D	10ms	4ms

- The **RR** time-slice is **2ms**.
 - Suppose when a job arrives, it is added to the tail of a work queue. The **RR** policy selects the next job of the current job in the queue.
1. Please calculate the average turnaround time and average response time for various scheduling policies. ($8 \times 1' = 8'$)

Job	Average Turnaround Time	Average Response Time
FIFO	[1]	[2]
SJF	[3]	[4]
STCF	[5]	[6]
RR	[7]	[8]

2. We decide to use **MLFQ** scheduling policy now. The workload remains the same.
- There are **2 priority queues**. The highest one has time-slice of **1 ms**, the lowest one has time-slice of **2ms**.
 - **RR** is used in each queue.
 - The CPU scheduling is carried out only at **completion** of processes or time-slices.

Following table shows the execution of CPU. No I/O issues are involved. Please fill in the blanks ($8 \times 1' = 8'$).

Time	0	1	2	3	4	5	6	7	8	9	10
CPU	A	A	A	A	[1]	[2]	[3]	B	B	C	[4]
Time	11	12	13	14	15	16	17	18	19	20	21
CPU	[5]	D	D	A	A	B	[6]	C	D	A	C

Average turnaround time: [7] ms

(HINT: **DON'T** forget the last executing time-slice when calculating turnaround time).

Average response time: [8] ms

题号: **W2**

Process (16 points)

Assume we have the following codes. We have following jobs in the workload. No I/O issues.

```
1. // headers are omitted
2.
3. int num = 3;
4. int main(void) {
5.     for (int i = 0; i < 2; ++i) {
6.         pid_t pid = fork();
7.         if (pid == 0) {
8.             num++;
9.             printf("child with i: %d, num: %d\n", i, num);
10.        } else {
11.            waitpid(pid, NULL, 0);
12.            printf("parent with i: %d, num: %d\n", i, num);
13.        }
14.        num <<= 1;
15.    }
16.
17.    return 0;
18. }
```

1. How many times will `fork()` be called in this program? How many times will `printf()` be called in this program($2*2' = 4'$)?
2. If `waitpid()` in line 11 is removed, will the times that `fork()` is called change ($2'$)?
3. Please give a possible output of the program. You can use [`<P/C>`,`<i>`,`<num>`] (e.g., [`C,0,4`]) to represent each printed line for simplicity. And you can write your answer in one line and neglect `'\n'` ($6'$).
4. Please explain constraints that `waitpid()` puts on the output of the program($4'$).

题号: X1

Signal1 (18 points)

```
1. #include <csapp.h>
2. #define MAX 2
3. #define SEC 4
4. int cnt = 0;
5. void handler(int n) {
6.     if (cnt < MAX) {
7.         cnt++;
8.         Alarm(2);
9.         Fork();
10.        printf("%d\n", cnt);
11.    }
12. }
13.
14. int main(void) {
15.     Signal(SIGALRM, handler);
16.     Alarm(2);
17.     for (int i = 0; i < SEC; i++) {
18.         Sleep(1);
19.         printf("S%d\n", i+1);
20.     }
21.     return 0;
22. }
```

One of the TAs wrote a program as shown above. NOTE:

- Assume all system calls are successful.
 - A child created via **fork** inherits a copy of its parent's set of currently blocked signals.
 - A child created via **fork** initially has an **empty** pending signal set.
 - **alarm** before **fork** will **only** send signal to parent process later.
 - We assume that the execution time of **sleep(1)** is strictly 1 second, the same for **alarm**.
1. Please write one possible output on the screen. (Note: you don't have to fill out all the blanks) (3')

Line1	Line2	Line3	Line4	Line5	Line6	Line7	Line8
S1							
Line9	Line10	Line11	Line12	Line13	Line14	Line15	Line16

--	--	--	--	--	--	--	--

2. Suppose **MAX** is **4** and **SEC** is **5** in this question.
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
3. Suppose **MAX** is **4** and **SEC** is **10** in this question.
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
4. Suppose we exchange the codes in **line 8** and **line 9** in this question (Note: **MAX** is **2** and **SEC** is **4**).
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
 - c) Is there any race modifying the global variable cnt between the parent process and its child process(es), why? (3')

题号: X2

Signal2 (18 points)

```
1. #include <csapp.h>
2. #define MAX 2
3. #define SEC 4
4. int cnt = MAX;
5. void handler(int n) {
6.     if (cnt > 0) {
7.         cnt--;
8.         Alarm(2);
9.         Fork();
10.        printf("%d\n", cnt);
11.    }
12. }
13.
14. int main(void) {
15.     Signal(SIGALRM, handler);
16.     Alarm(2);
17.     for (int i = 0; i < SEC; i++) {
18.         Sleep(1);
19.         printf("S%d\n", i+1);
20.     }
21.     return 0;
22. }
```

One of the TAs wrote a program as shown above. NOTE:

- Assume all system calls are successful.
- A child created via **fork** inherits a copy of its parent's set of currently blocked signals.
- A child created via **fork** initially has an **empty** pending signal set.
- **alarm** before **fork** will **only** send signal to parent process later.
- We assume that the execution time of **sleep(1)** is strictly 1 second, the same for **alarm**.

1. Please write one possible output on the screen. (Note: you don't have to fill out all the blanks) (3')

Line1	Line2	Line3	Line4	Line5	Line6	Line7	Line8
S1							
Line9	Line10	Line11	Line12	Line13	Line14	Line15	Line16

--	--	--	--	--	--	--	--

2. Suppose **MAX** is **4** and **SEC** is **6** in this question.
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
3. Suppose **MAX** is **5** and **SEC** is **9** in this question.
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
4. Suppose we exchange the codes in **line 8** and **line 9** in this question (Note: **MAX** is **2** and **SEC** is **4**).
 - a) How many lines of output are on the screen when this program ends? (2')
 - b) How many processes are there before the end of the program? (2')
 - c) Is there any race modifying the global variable cnt between the parent process and its child process(es), why? (3')

题号: Y1

Network1 (24 points)

Your task is to design a simple system to get your lab grades. This system consists of one server and several clients. The hostname of the server is `ipads.se.sjtu.edu.cn`, and the port number is 8080.

Suppose you want to query your grades of one of the labs. Firstly, you need to send a request to get the lab lists that you have handed in. Then you will know which lab's grade you can query about. Finally, you need to send a request again to get the lab's grade.

The detailed communication protocol between the server and the clients is defined as follows:

1. A client sends the **LABHISTORY** request to get the submission history of student with id `"518xxxxxxxx"`. The password of this student must be provided, and is simply transferred in plaintext:

```
"LABHISTORY" '\n' "518xxxxxxxx" '\n' "myPassword" '\n'
```

2. The server receives the **LABHISTORY** request and returns the following messages containing the submission history of this student:

```
"lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' ... "lab10" '\n' "END  
LABHISTORY" '\n'
```

If the password is incorrect or the student id does not exist, the server simply returns error message:

```
"INCORRECT ACCOUNT\n"
```

3. If the server returns error message, or the name of the lab provided by the user does not appear in the lab history list, then client prints corresponding error message:

```
"Incorrect account!\n", or
```

```
"Invalid labname!\n"
```

4. If the name of the lab is found in the lab history list, the client will send another **GRADE** request to get the grade:

```
"GRADE" '\n' "lab3" '\n' "518xxxxxxxx" '\n' "myPassword" '\n'
```

5. The server receives the **GRADE** request and returns the corresponding grade:

```
"95\n"
```

We do not consider incorrect account in this step, since it's already checked.

Examples:

1] Command: ./getGrade lab3 51800001111 myPassword

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "myPassword" '\n'

Server => Client, sends response: "lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' "END LABHISTORY" '\n'

Client => Server, sends request: "GRADE" '\n' "lab3" '\n' "51800001111" '\n' "myPassword" '\n'

Server => Client, sends response: "95\n"

Client prints "95\n"

2] Command: ./getGrade lab6 51800001111 myPassword

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "myPassword" '\n'

Server => Client, sends response: "lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' "END LABHISTORY" '\n'

Client prints "Invalid labname!\n"

3] Command: ./getGrade lab3 51800001111 wrongPass

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "wrongPass" '\n'

Server => Client, sends response: "INCORRECT ACCOUNT\n"

Client prints "Incorrect account!\n"

Please complete the following client-side code. **Hint:** you can use `sprintf` to construct the requests. Also you can use `sscanf` to parse the messages from the server.

```
#include "csapp.h"

/* the program is executed by ./queryInfo <labname> <studentid> <password> */

int main(int argc, char **argv) {

    char *labname = argv[1];

    char *stu_id  = argv[2];

    char *pass    = argv[3];

    char send_buf[1024], recv_buf[1024];

    memset(send_buf, 0, 1024);

    memset(recv_buf, 0, 1024);


    /* connect to <hostname>:<port> and send out the request */

    [1] // Write your code here (5')


    /* read the labhistory information from server */

    /* check the lab's existence */

    [2] // Write your code here (9')


    /* send GRADE request */

    [3] // Write your code here (5')
```

```
/* finally get the grade information
```

```
and print to the screen */
```

```
[4] // Write your code here (5')
```

```
return 0;
```

```
}
```

题号: Y2

Network (24 points)

Your task is to design a simple system to get your lab grades. This system consists of one server and several clients. The hostname of the server is `ipads.se.sjtu.edu.cn`, and the port number is `8888`.

Suppose you want to query your grades of one of the labs. Firstly, you need to send a request to get the lab lists that you have handed in. Then you will know which lab's grade you can query about. Finally, you need to send a request again to get the lab's grade.

The detailed communication protocol between the server and the clients is defined as follows:

1. A client sends the **LABHISTORY** request to get the submission history of student with id `"518xxxxxxxx"`. The password of this student must be provided, and is simply transferred in plaintext:

```
"LABHISTORY" '\n' "518xxxxxxxx" '\n' "myPassword" '\n'
```

2. The server receives the **LABHISTORY** request and returns the following messages containing the submission history of this student:

```
"LABHISTORY BEGIN" '\n' "lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' ...  
"lab10" '\n' "LABHISTORY END" '\n'
```

If the password is incorrect or the student id does not exist, the server simply returns error message:

```
"INCORRECT ACCOUNT\n"
```

3. If the server returns error message, or the name of the lab provided by the user does not appear in the lab history list, then client prints corresponding error message:

```
"Incorrect account!\n", or
```

```
"Invalid labname!\n"
```

4. If the name of the lab is found in the lab history list, the client will send another **GRADE** request to get the grade:

```
"GRADE" '\n' "518xxxxxxxx" '\n' "myPassword" '\n' "lab3" '\n'
```

5. The server receives the **GRADE** request and returns the corresponding grade:

```
"95\n"
```

We do not consider incorrect account in this step, since it's already checked.

Examples:

1] Command: ./getGrade 51800001111 myPassword lab3

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "myPassword" '\n'

Server => Client, sends response: "LABHISTORY BEGIN" '\n' "lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' "LABHISTORY END" '\n'

Client => Server, sends request: "GRADE" '\n' "51800001111" '\n' "myPassword" '\n' "lab3" '\n'

Server => Client, sends response: "95\n"

Client prints "95\n"

2] Command: ./getGrade 51800001111 myPassword lab6

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "myPassword" '\n'

Server => Client, sends response: "LABHISTORY BEGIN" '\n' "lab1" '\n' "lab2" '\n' "lab3" '\n' "lab4" '\n' "LABHISTORY END" '\n'

Client prints "Invalid labname!\n"

3] Command: ./getGrade 51800001111 wrongPass lab3

Client => Server, sends request: "LABHISTORY" '\n' "51800001111" '\n' "wrongPass" '\n'

Server => Client, sends response: "INCORRECT ACCOUNT\n"

Client prints "Incorrect account!\n"

Please complete the following client-side code. **Hint:** you can use `sprintf` to construct the requests. Also you can use `sscanf` to parse the messages from the server.

```
#include "csapp.h"

/* the program is executed by ./queryInfo <labname> <studentid> <password> */

int main(int argc, char **argv) {

    char *stu_id  = argv[1];

    char *pass    = argv[2];

    char *labname = argv[3];

    char send_buf[1024], recv_buf[1024];

    memset(send_buf, 0, 1024);

    memset(recv_buf, 0, 1024);

    /* connect to <hostname>:<port> and send out the request */

    [1] // Write your code here (5')

    /* read the labhistory information from server */

    /* check the lab's existence */

    [2] // Write your code here (9')

    /* send GRADE request */

    [3] // Write your code here (5')
```

```
/* finally get the grade information
```

```
and print to the screen */
```

```
[4] // Write your code here (5')
```

```
return 0;
```

```
}
```

题号: **Z1**

Lock1 (26 points)

Sam modifies the spin lock by adding one line "**IDLE**(...);" in the while-loop.

- **IDLE**(*n*) will consume $C*n$ CPU cycles where *C* is a user-defined **constant**.
- **TAS** will **atomically test-and-set** the content of the target address
- When there are **multiple** threads **TAS** the same address concurrently, it will take **much longer** to finish this operation

The original spin lock	A modified spin lock
<pre>1. typedef struct __lock_t { 2. int flag; 3. } lock_t; // init to 0 4. void lock(lock_t *lock) { 5. while (TAS(&lock->flag, 1)) 6. ; 7. } 8. void unlock(lock_t *lock) { 9. lock->flag = 0; 10.}</pre>	<pre>1. typedef struct __lock_t { 2. int flag; 3. } lock_t; // init to 0 4. void lock(lock_t *lock) { 5. while (TAS(&lock->flag, 1)) 6. IDLE(...); 7. } 8. void unlock(lock_t *lock) { 9. lock->flag = 0; 10.}</pre>

1. Does this lock work **correctly**? Why or why not? (3')
2. Can this lock guarantee the acquire **fairness**? Explain your answer through an example. (3')
3. What's the **advantage** after adding the "**IDLE**(...);" line? (3')
4. What's the **disadvantage** after adding the "**IDLE**(...);" line? (3')
HINT: Consider what if an inappropriate constant *C* is chosen.

5. Since **IDLE** still consume CPU time, can we use **park** instead? The modified version of spin lock is shown below. Will this lock work correctly? Explain your answer through an example. (**FAA** will **atomically fetch-and-add** the target address with the given value) (6')

```
1. typedef struct __lock_t { int flag; int wait; } lock_t;
2. void lock(lock_t *lock) {
3.     while (TAS(&lock->flag, 1)) {
4.         FAA(&lock->wait, 1);
5.         park();
6.     }
7. }
8. void unlock(lock_t *lock) {
9.     lock->flag = 0;
10.    if (lock->wait) {
11.        FAA(&lock->wait, -1);
12.        unpark();
13.    }
14.}
```

6. Rather than **IDLE** for a certain time, can you come up with some better idea? Write down you own version of spin lock. (6')

(HINT: Since executing **TAS** concurrently will cause significant overhead, you can use **while** to wait on some condition before executing **TAS** again.)

题号: **Z2**

Lock2 (26 points)

Sam modifies the ticket lock by adding one line "**IDLE(...);**" in the while-loop.

- **IDLE(n)** will consume $C*n$ CPU cycles where C is a user-defined **constant**.
- **FAA** will **atomically fetch-and-add** the content of the target address
- When there are **multiple** threads access the same address concurrently, it will take **much longer** to finish the memory operation

The original ticket lock	A modified ticket lock
<pre>typedef struct __lock_t { int ticket; int turn; } lock_t; void lock(lock_t *lock) { int myturn = FAA(&lock->ticket); while (lock->turn != myturn) ; } void unlock(lock_t *lock) { lock->turn ++; }</pre>	<pre>typedef struct __lock_t { int ticket; int turn; } lock_t; void lock(lock_t *lock) { int myturn = FAA(&lock->ticket); while (lock->turn != myturn) IDLE(...); } void unlock(lock_t *lock) { lock->turn ++; }</pre>

1. Does this lock work **correctly**? Why or why not? (3')
2. Can this lock guarantee the acquire **fairness**? Why? (3')
3. What's the **advantage** after adding the "**IDLE(...);**" line? (3')
4. What's the **disadvantage** after adding the "**IDLE(...);**" line? (3')
HINT: Consider what if an inappropriate constant C is chosen.

5. Since **IDLE** still consume CPU time, can we use **park** instead? The modified version of spin lock is shown below. Will this lock work correctly? Explain your answer through an example. (**FAA** will **atomically fetch-and-add** the target address with the given value) (6')

```
1. typedef struct __lock_t { int ticket; int turn; } lock_t;
2. void lock(lock_t *lock) {
3.     int myturn = FAA(&lock->ticket);
4.     while (lock->turn != myturn) {
5.         park();
6.     }
7. }
8. void unlock(lock_t *lock) {
9.     lock->turn ++;
10.    if (lock->ticket > lock->turn) {
11.        unpark();
12.    }
13.}
```

6. Rather than asking all threads to **IDLE** for a same number of cycles, can you come up with some better idea? Write down you own version of ticket lock. (6')